

```

In [57]: #2.3 Missing values
import pandas as pd
df = pd.read_csv(r"C:\Users\vaibh\Downloads\heart (1).csv")
df.head(10)

df_dropped = df.dropna()
df_dropped

#filling with specific value
df_specific = df.fillna(0)
df_specific

df_mean = df.fillna(df.mean())
df_mean

```

```

Out[57]:

```

	age	gender	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope
0	63.000000	1.000000	3	145.0	233	1	0	150	0	2.3	0
1	54.413333	0.684385	2	130.0	250	0	1	187	0	3.5	0
2	41.000000	0.000000	1	130.0	204	0	0	172	0	1.4	2
3	54.413333	1.000000	1	120.0	236	0	1	178	0	0.8	2
4	57.000000	0.684385	0	120.0	354	0	1	163	1	0.6	2
...
298	57.000000	0.000000	0	140.0	241	0	1	123	1	0.2	1
299	45.000000	1.000000	3	110.0	264	0	1	132	0	1.2	1
300	68.000000	1.000000	0	144.0	193	1	1	141	0	3.4	1
301	57.000000	1.000000	0	130.0	131	0	1	115	1	1.2	1
302	57.000000	0.000000	1	130.0	236	0	0	174	0	0.0	1

303 rows × 14 columns



```

In [134... # 2.3 remove outliers
import pandas as pd
import numpy as np
from scipy.stats import zscore

df = pd.read_csv(r"C:\Users\vaibh\Downloads\heart (1).csv")

z_score = zscore(df, nan_policy='omit')
abs_z_score = np.abs(z_score)
threshold = 2
abs_z_score.head(100)

```

Out[134...

	age	gender	cp	trestbps	chol	fbs	restecg	thalach	exan
0	0.948186	0.679091	1.973123	0.764565	0.256334	2.394438	1.005832	0.015443	0.69663
1	NaN	NaN	1.002577	0.091038	0.072199	0.417635	0.898962	1.633471	0.69663
2	1.481172	1.472556	0.032031	0.091038	0.816773	0.417635	1.005832	0.977514	0.69663
3	NaN	0.679091	0.032031	0.661440	0.198357	0.417635	0.898962	1.239897	0.69663
4	0.285634	NaN	0.938515	0.661440	2.082050	0.417635	0.898962	0.583939	1.43548
...
95	0.156068	0.679091	0.938515	0.593445	0.391612	0.417635	1.005832	1.690047	1.43548
96	0.837760	1.472556	0.938515	0.479364	2.855069	0.417635	1.005832	0.321556	0.69663
97	0.266493	0.679091	0.938515	1.345922	0.256334	2.394438	0.898962	0.115749	0.69663
98	1.260321	0.679091	1.002577	0.091038	1.328356	0.417635	0.898962	0.540209	0.69663
99	0.156068	0.679091	1.002577	0.091038	0.005102	2.394438	1.005832	1.021244	0.69663

100 rows × 14 columns



In [142...

```
df_cleaned = df.loc[(abs_z_score < threshold).all(axis=1)]  
df_cleaned.head(100)
```

Out[142...

	age	gender	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	t
2	41.0	0.0	1	130.0	204	0	0	172	0	1.4	2	0	
7	44.0	1.0	1	120.0	263	0	1	173	0	0.0	2	0	
9	57.0	1.0	2	150.0	168	0	1	174	0	1.6	2	0	
10	54.0	1.0	0	140.0	239	0	1	160	0	1.2	2	0	
11	48.0	0.0	2	130.0	275	0	1	139	0	0.2	2	0	
...
146	44.0	0.0	2	118.0	242	0	1	149	0	0.3	1	1	
147	60.0	0.0	3	150.0	240	0	1	171	0	0.9	2	0	
148	44.0	1.0	2	120.0	226	0	1	169	0	0.0	2	0	
149	42.0	1.0	2	130.0	180	0	1	150	0	0.0	2	0	
151	71.0	0.0	0	112.0	149	0	1	125	0	1.6	1	0	

100 rows × 14 columns



In [150...

```
#Api integration
import pandas as pd
import requests

url = "https://api.coingecko.com/api/v3/coins/markets"

params = {
    "vs_currency" : "usd",
    "order" : "market_cap_desc",
    "per_page" : 10,
    "page" : 1,
    "sparkLine" : False
}

response = requests.get(url,params=params)

if response.status_code == 200 :
    data = response.json()
    df = pd.DataFrame(data)
    print("Data succesfully loaded into dataframe")
else :
    print(f"Failed to fetch the data : {response.status_code}")

df.head()
```

Data succesfully loaded into dataframe

Out[150...

	id	symbol	name	image	current_price	ma
0	bitcoin	btc	Bitcoin	https://coin-images.coingecko.com/coins/images/1/bitcoin.png	92380.000000	182755
1	ethereum	eth	Ethereum	https://coin-images.coingecko.com/coins/images/279/ethereum.png	3113.340000	37480
2	tether	usdt	Tether	https://coin-images.coingecko.com/coins/images/320/tether.png	0.999951	12811
3	solana	sol	Solana	https://coin-images.coingecko.com/coins/images/256/solana.png	241.670000	11468
4	binancecoin	bnb	BNB	https://coin-images.coingecko.com/coins/images/83/binance-coin.png	615.400000	8991

5 rows × 26 columns



In [152...

```
Out[152... Index(['id', 'symbol', 'name', 'image', 'current_price', 'market_cap',
      'market_cap_rank', 'fully_diluted_valuation', 'total_volume',
      'high_24h', 'low_24h', 'price_change_24h',
      'price_change_percentage_24h', 'market_cap_change_24h',
      'market_cap_change_percentage_24h', 'circulating_supply',
      'total_supply', 'max_supply', 'ath', 'ath_change_percentage',
      'ath_date', 'atl', 'atl_change_percentage', 'atl_date', 'roi',
      'last_updated'],
      dtype='object')
```

```
In [185... print("Column names in DataFrame : ")
print(df.columns)

df_selected = df[['name', 'current_price', 'market_cap']]
print("\n Selected Columns : ")
print(df_selected.head())

df_sorted = df_selected.sort_values(by='market_cap', ascending = False)
print("\nTop Cryptocurrencies of by market cap : ")
print(df_sorted)

df_selected['price_in_eur'] = df_selected['current_price'] * 0.85
print("\n Added 'price in eur' column : ")
print(df_selected.head())

df_filtered = df_selected[df_selected['market_cap'] > 1e9]
print("Cryptocurrencies with market cap greter than 1 billion")
print(df_filtered)
```

Column names in DataFrame :

```
Index(['id', 'symbol', 'name', 'image', 'current_price', 'market_cap',  
      'market_cap_rank', 'fully_diluted_valuation', 'total_volume',  
      'high_24h', 'low_24h', 'price_change_24h',  
      'price_change_percentage_24h', 'market_cap_change_24h',  
      'market_cap_change_percentage_24h', 'circulating_supply',  
      'total_supply', 'max_supply', 'ath', 'ath_change_percentage',  
      'ath_date', 'atl', 'atl_change_percentage', 'atl_date', 'roi',  
      'last_updated'],  
      dtype='object')
```

Selectd Columns :

	name	current_price	market_cap
0	Bitcoin	92380.000000	1827556947075
1	Ethereum	3113.340000	374806136312
2	Tether	0.999951	128114926302
3	Solana	241.670000	114682944570
4	BNB	615.400000	89913267776

Top Cryptocurrencies of by market cap :

	name	current_price	market_cap
0	Bitcoin	92380.000000	1827556947075
1	Ethereum	3113.340000	374806136312
2	Tether	0.999951	128114926302
3	Solana	241.670000	114682944570
4	BNB	615.400000	89913267776
5	XRP	1.097000	62439424379
6	Dogecoin	0.396368	58184329728
7	USDC	0.999439	37239890081
8	Lido Staked Ether	3112.840000	30426780421
9	Cardano	0.733744	26307829353

Added 'price in eur' column :

	name	current_price	market_cap	price_in_eur
0	Bitcoin	92380.000000	1827556947075	78523.000000
1	Ethereum	3113.340000	374806136312	2646.339000
2	Tether	0.999951	128114926302	0.849958
3	Solana	241.670000	114682944570	205.419500
4	BNB	615.400000	89913267776	523.090000

Cryptocurrencies with market cap greter than 1 billion

	name	current_price	market_cap	price_in_eur
0	Bitcoin	92380.000000	1827556947075	78523.000000
1	Ethereum	3113.340000	374806136312	2646.339000
2	Tether	0.999951	128114926302	0.849958
3	Solana	241.670000	114682944570	205.419500
4	BNB	615.400000	89913267776	523.090000
5	XRP	1.097000	62439424379	0.932450
6	Dogecoin	0.396368	58184329728	0.336913
7	USDC	0.999439	37239890081	0.849523
8	Lido Staked Ether	3112.840000	30426780421	2645.914000
9	Cardano	0.733744	26307829353	0.623682

C:\Users\vaibh\AppData\Local\Temp\ipykernel_11488\1469776383.py:12: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_selected['price_in_eur'] = df_selected['current_price'] * 0.85
```

In [201...

```
import sqlite3

connection = sqlite3.connect("sample_database.db")
cursor = connection.cursor()

#create table
cursor.execute("""
CREATE TABLE IF NOT EXISTS employees (
    id INTEGER PRIMARY KEY AUTOINCREMENT,
    name TEXT NOT NULL,
    age INTEGER,
    department TEXT,
    salary REAL

)
""")
cursor.executemany("""
INSERT INTO employees (name,age,department,salary)
VALUES(?,?,?,?)
""", [
    ("Allice",30,"HR",600000),
    ("Bob",25,"IT",70000),
    ("Charlie", 35, "Finance", 80000),
    ("Diana", 28, "IT", 65000),
    ("Eve", 40, "HR", 72000)

])
connection.commit()
connection.close()
print("Data and table setup complete")
```

Data and table setup complete

In [203...

```
import pandas as pd
connection = sqlite3.connect("sample_database.db")

query = "SELECT * FROM employees"
df = pd.read_sql_query(query,connection)

print("Data loaded into Dataframe : ")
print(df)

connection.close()
```

Data loaded into Dataframe :

	id	name	age	department	salary
0	1	Alice	30	HR	600000.0
1	2	Bob	25	IT	70000.0
2	3	Alice	30	HR	600000.0
3	4	Bob	25	IT	70000.0
4	5	Charlie	35	Finance	80000.0
5	6	Diana	28	IT	65000.0
6	7	Eve	40	HR	72000.0

In [215...

```
print("\n Summary Statistics")
print(df.describe())

df["tax"] = df["salary"]*0.10
print("\n DataFrame with Tax column : ")
print(df)

high_earners = df[df["salary"]>65000]
print("\n Employees with salary > 65000")
print(high_earners)

avg_sal_by_dept = df.groupby("department")["salary"].mean()
print("\nAverage salary by the department : ")
print(avg_sal_by_dept)

sorted_by_age = df.sort_values(by="age",ascending = False)
print("\nEmployees Sorted Age : ")
print(sorted_by_age)
```

Summary Statistics

	id	age	salary	tax
count	7.000000	7.000000	7.000000	7.000000
mean	4.000000	30.428571	222428.571429	22242.857143
std	2.160247	5.442338	257968.898088	25796.889809
min	1.000000	25.000000	65000.000000	6500.000000
25%	2.500000	26.500000	70000.000000	7000.000000
50%	4.000000	30.000000	72000.000000	7200.000000
75%	5.500000	32.500000	340000.000000	34000.000000
max	7.000000	40.000000	600000.000000	60000.000000

DataFrame with Tax column :

	id	name	age	department	salary	tax
0	1	Allice	30	HR	600000.0	60000.0
1	2	Bob	25	IT	70000.0	7000.0
2	3	Allice	30	HR	600000.0	60000.0
3	4	Bob	25	IT	70000.0	7000.0
4	5	Charlie	35	Finance	80000.0	8000.0
5	6	Diana	28	IT	65000.0	6500.0
6	7	Eve	40	HR	72000.0	7200.0

Employees with salary > 65000

	id	name	age	department	salary	tax
0	1	Allice	30	HR	600000.0	60000.0
1	2	Bob	25	IT	70000.0	7000.0
2	3	Allice	30	HR	600000.0	60000.0
3	4	Bob	25	IT	70000.0	7000.0
4	5	Charlie	35	Finance	80000.0	8000.0
6	7	Eve	40	HR	72000.0	7200.0

Average salary by the department :

department

Finance 80000.000000

HR 424000.000000

IT 68333.333333

Name: salary, dtype: float64

Employees Sorted Age :

	id	name	age	department	salary	tax
6	7	Eve	40	HR	72000.0	7200.0
4	5	Charlie	35	Finance	80000.0	8000.0
0	1	Allice	30	HR	600000.0	60000.0
2	3	Allice	30	HR	600000.0	60000.0
5	6	Diana	28	IT	65000.0	6500.0
1	2	Bob	25	IT	70000.0	7000.0
3	4	Bob	25	IT	70000.0	7000.0

```
In [77]: import pandas as pd
import requests
from bs4 import BeautifulSoup
from io import StringIO

url = "https://en.wikipedia.org/wiki/Agriculture_in_India"

response = requests.get(url)
```



```
soup = BeautifulSoup(response.text,"html.parser")
table = soup.find("table",{ "class" : "wikitable"})

df = pd.read_html(StringIO(str(table)))[0]

print("Extracted Data : ")
df.head()
```

Extracted Data :

Out[77]:

	Rank	Commodity	Value (US\$, 2016)	Unit price (US\$ / kilogram, 2009)	Average yield (tonnes per hectare, 2017)	Most productive country (tonnes per hectare, 2017)	Most productive country (tonnes per hectare, 2017).1
0	1	Rice	\$70.18 billion	0.27	3.85	9.82	Australia
1	2	Buffalo milk	\$43.09 billion	0.40	2.00[78]	2.00[78]	India
2	3	Cow milk	\$32.55 billion	0.31	1.2[78]	10.3[78]	Israel
3	4	Wheat	\$26.06 billion	0.15	2.8	8.9	Netherlands
4	5	Cotton (Lint + Seeds)	\$23.30 billion	1.43	1.6	4.6	Israel

```
In [52]: import pandas as pd

data = {
    'name' : ["Vaibhav","Riya","Taniya","Tanmay","Yashraj"],
    'Age' : [12,14,16,17,18],
    'rank' : [1,3,2,5,4],
    'score' : [99,67,87,12,34]
}
df = pd.DataFrame(data)
print(df)
print(df.columns)
df_selected = df[df['Age'] > 14]
print(df_selected)
df_sorted = df.sort_values(by="rank",ascending = False)
print(df_sorted)
df_grouping = df.groupby("name")["score"].mean()
print(df_grouping)
```

	name	Age	rank	score
0	Vaibhav	12	1	99
1	Riya	14	3	67
2	Taniya	16	2	87
3	Tanmay	17	5	12
4	Yashraj	18	4	34

Index(['name', 'Age', 'rank', 'score'], dtype='object')

	name	Age	rank	score
2	Taniya	16	2	87
3	Tanmay	17	5	12
4	Yashraj	18	4	34

	name	Age	rank	score
3	Tanmay	17	5	12
4	Yashraj	18	4	34
1	Riya	14	3	67
2	Taniya	16	2	87
0	Vaibhav	12	1	99

name

Riya 67.0

Taniya 87.0

Tanmay 12.0

Vaibhav 99.0



Yashraj 34.0

Name: score, dtype: float64

```
In [87]: import numpy as np
import matplotlib.pyplot as plt
A = np.arange(1,11)
B = np.random.randint(1,20,size = 10)

plt.subplot(1, 2, 1) # 1 row, 2 columns, first subplot
plt.plot(A,B,label="Lineplot",marker = "o",linestyle = "--",color = "blue")
plt.xlabel("X-axis")
plt.ylabel("Y-axis")
plt.legend()
plt.title("LinePlot")
plt.grid(True)
plt.ylim(0, 25) # Set y-axis limits
plt.xticks(rotation=45) # Rotate x-axis labels
plt.show()

plt.subplot(1, 2, 1) # 1 row, 2 columns, first subplot
plt.bar(A,B,label = "BarPlot",color = "skyblue")
plt.xlabel("X-axis")
plt.ylabel("Y-label")
plt.legend()
plt.title("BarPlot")
plt.grid(axis="y") # Grid only on the y-axis
plt.ylim(0, 25) # Set y-axis limits
plt.xticks(rotation=45) # Rotate x-axis labels
plt.show()
```



In [106...

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv(r"C:\Users\vaibh\Downloads\iris.csv")
print(df.describe())
print()
print(df.info())
print(df.isnull().sum())
```

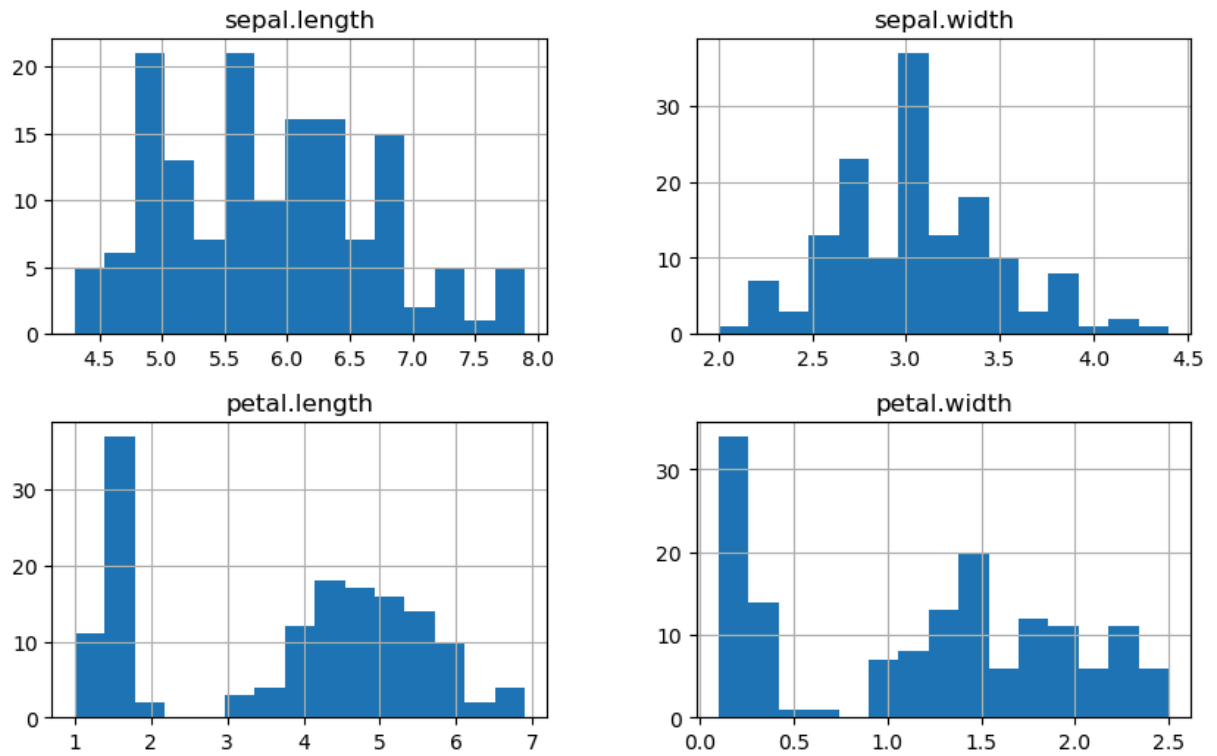
	sepal.length	sepal.width	petal.length	petal.width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   sepal.length    150 non-null   float64
1   sepal.width     150 non-null   float64
2   petal.length    150 non-null   float64
3   petal.width     150 non-null   float64
4   species         150 non-null   object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
None
sepal.length    0
sepal.width     0
petal.length    0
petal.width     0
species         0
dtype: int64
```


In [114...

```
df[['sepal.length', 'sepal.width', 'petal.length', 'petal.width']].hist(bins = 15, figs
plt.suptitle("Features Histigram")
plt.show()
```

Features Histogram

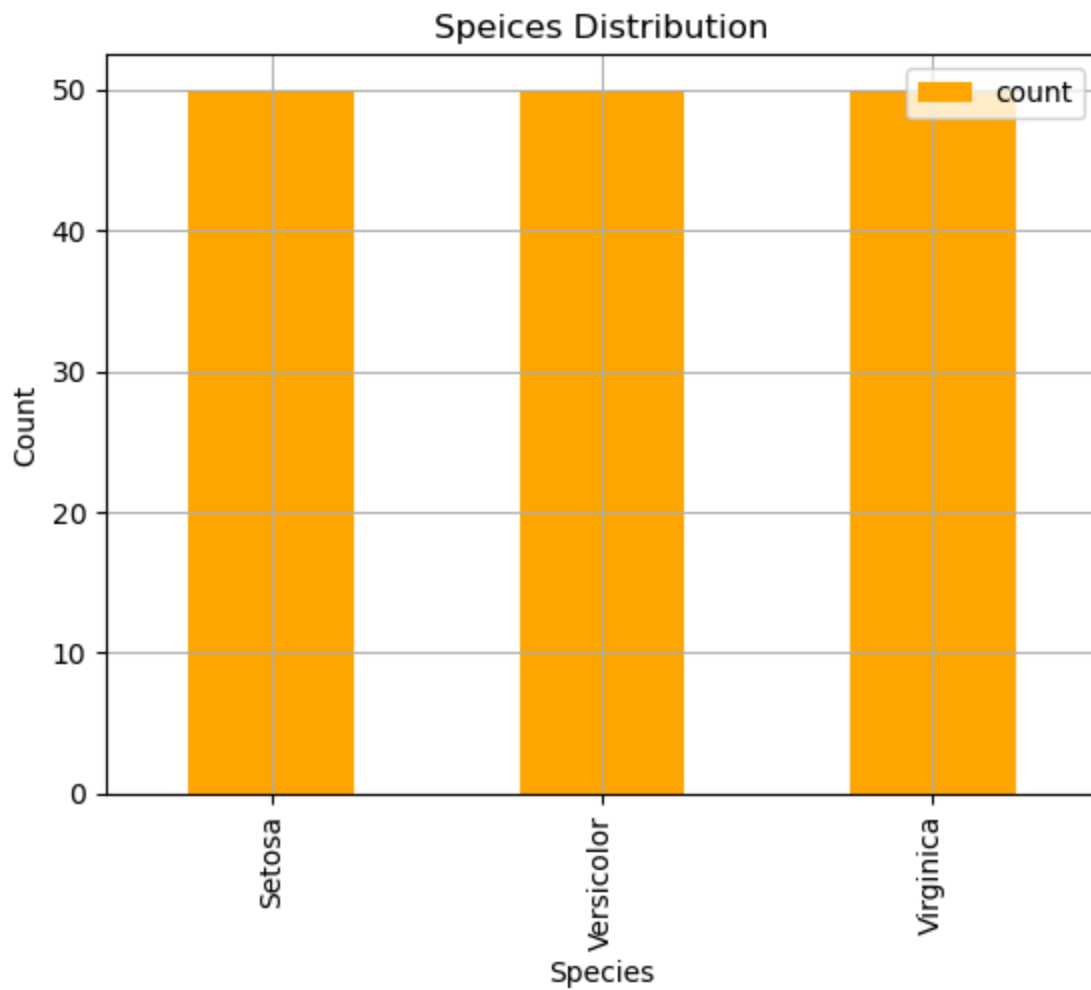


```
In [122... plt.scatter(df['sepal.length'],df['petal.length'],color = 'blue',alpha = 0.5)
plt.title("Scatter plot Sepal.length vs Petal.length")
plt.xlabel("Sepal.length")
plt.ylabel("Petal.length")
plt.show()
```




In [132...

```
df['species'].value_counts().plot(kind = 'bar',color = 'orange')
plt.title("Speices Distribution")
plt.xlabel("Species")
plt.ylabel("Count")
plt.grid(True)
plt.legend()
plt.show()
```



In [140...

```
import seaborn as sns
corelation_matrix = df.select_dtypes(include = 'number').corr()
plt.figure(figsize = (10,6))
sns.heatmap(corelation_matrix,annot = True,cmap = 'coolwarm',fmt='0.2f',cbar = True)
plt.show()
```



```
In [3]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
df = pd.read_csv(r"C:\Users\vaibh\Downloads\Orange_Telecom_Churn_Data.csv")
df.head(5)
```

```
Out[3]:
```

	state	account_length	area_code	phone_number	intl_plan	voice_mail_plan	number_vm
0	KS	128	415	382-4657	no	yes	
1	OH	107	415	371-7191	no	yes	
2	NJ	137	415	358-1921	no	no	
3	OH	84	408	375-9999	yes	no	
4	OK	75	415	330-6626	yes	no	

5 rows × 7 columns



```
In [9]: sns.histplot(df['total_day_charge'], kde = True, color = 'blue')
plt.title("Histogram of total day charge")
plt.xlabel("Total Day Charge")
plt.ylabel("Frequency")
plt.grid(True)
plt.show()

sns.boxplot(data = df[['total_day_charge', 'total_eve_calls']])
plt.title("Total Day Charge vs Total eve calls")
plt.xlabel("Total Day Charge")
```




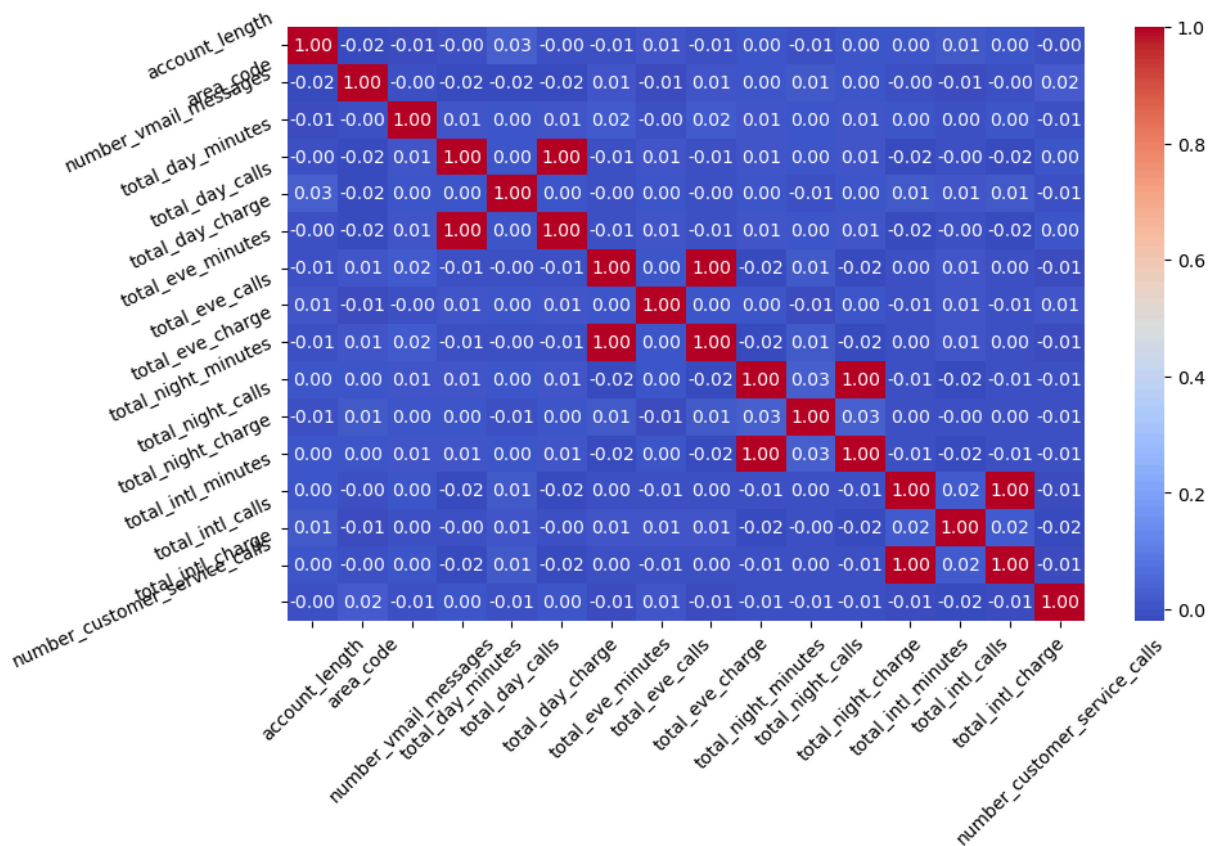
```

plt.ylabel("Total eve cells")
plt.xticks(rotation = 45)
plt.grid(True)
plt.show()

correlation_matrix = df.select_dtypes(include = 'number').corr()
plt.figure(figsize = (10,6))
sns.heatmap(correlation_matrix,annot= True,cmap = 'coolwarm',fmt = '0.2f',cbar = Tr
plt.xticks(rotation = 45)
plt.yticks(rotation = 25)
plt.show()

```






```
In [129... import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
df = pd.read_csv(r"C:\Users\vaibh\Downloads\iris.csv")
print("Data frame with missing values : ")
print(df.isnull().sum())
df.dropna(inplace = True)
print("\n Number of Duplicates rows : ")
print(df.duplicated().sum())
df.drop_duplicates(inplace = True)
correalation_matrix = df.select_dtypes(include = 'number').corr()
print("\n Data set after Cleaning ")
print(df.info())
plt.figure(figsize =(10,6))
sns.heatmap(correalation_matrix,annot =False,cmap='coolwarm',fmt ='0.2f',cbar = True)
plt.show()
```

```
Data frame with missing values :
sepal.length    0
sepal.width     0
petal.length    0
petal.width     0
species         0
dtype: int64
```

```
Number of Duplicates rows :
1
```

```
Data set after Cleaning
<class 'pandas.core.frame.DataFrame'>
Index: 149 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   sepal.length    149 non-null    float64
1   sepal.width     149 non-null    float64
2   petal.length    149 non-null    float64
3   petal.width     149 non-null    float64
4   species         149 non-null    object
dtypes: float64(4), object(1)
memory usage: 7.0+ KB
None
```



```
In [190... import numpy as np
from scipy.stats import ttest_1samp

sample = [23, 27, 26, 30, 29, 22, 24, 28, 25, 31]
population_mean = 25
```

```

t_stat,p_value = ttest_1samp(sample,population_mean)
alpha = 0.05
print("T-statics : ",t_stat)
print("p_value : ",p_value)

if p_value < alpha :
    print("Reject the null hypothesis : The sample mean is completely difrent from p
else :
    print("Failed to reject the null hupothesis : there is no significant difference

```

T-statics : 1.5666989036012806

p_value : 0.1516274744876827

Failed to reject the null hupothesis : there is no significant difference between sam
ple mean and population mean

In [232...

```

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score,classification_report,confusion_matrix
import pandas as pd
df = pd.read_csv(r"C:\Users\vaibh\Downloads\heart (1).csv")
print("Missing values : ")
print(df.isnull().sum())
df.dropna(inplace = True)
print("\n Duplicated items : ")
print(df.duplicated().sum())
df.drop_duplicates(inplace = True)

X = df.iloc[:,0:13]
y = df['target']
print(X)
print(y)

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.3,random_state =
model = LogisticRegression(max_iter=1000)
model.fit(X_train,y_train)
y_pred = model.predict(X_test)

confusion = confusion_matrix(y_test,y_pred)
print("Confusion matrix ")
print(confusion)

acuuracy = accuracy_score(y_test,y_pred)
print("\nAccuracy Score : ")
print(acuuracy)

classification = classification_report(y_test,y_pred)
print("Classification Report : ")
print(classification)

plt.figure(figsize = (8,4))
sns.heatmap(df.corr(),annot = True,cmap = 'coolwarm',fmt = '0.2f',cbar = True)
plt.title("Corelation Matrix")
plt.show()

```

Missing values :

age 3
gender 2
cp 0
trestbps 1
chol 0
fbs 0
restecg 0
thalach 0
exang 0
oldpeak 0
slope 0
ca 0
thal 0
target 0
dtype: int64

Duplicated items :

1

	age	gender	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63.0	1.0	3	145.0	233	1	0	150	0	2.3	
2	41.0	0.0	1	130.0	204	0	0	172	0	1.4	
5	57.0	1.0	0	140.0	192	0	1	148	0	0.4	
7	44.0	1.0	1	120.0	263	0	1	173	0	0.0	
8	52.0	1.0	2	172.0	199	1	1	162	0	0.5	
..	
298	57.0	0.0	0	140.0	241	0	1	123	1	0.2	
299	45.0	1.0	3	110.0	264	0	1	132	0	1.2	
300	68.0	1.0	0	144.0	193	1	1	141	0	3.4	
301	57.0	1.0	0	130.0	131	0	1	115	1	1.2	
302	57.0	0.0	1	130.0	236	0	0	174	0	0.0	

	slope	ca	thal
0	0	0	1
2	2	0	2
5	1	0	1
7	2	0	3
8	2	0	3
..
298	1	0	3
299	1	0	3
300	1	2	3
301	1	1	3
302	1	1	2

[298 rows x 13 columns]

0 1
2 1
5 1
7 1
8 1
..
298 0
299 0
300 0
301 0

302 0

Name: target, Length: 298, dtype: int64

Confusion matrix


```
[[27  9]
 [ 6 48]]
```

Accuracy Score :

0.8333333333333334

Classification Report :

	precision	recall	f1-score	support
0	0.82	0.75	0.78	36
1	0.84	0.89	0.86	54
accuracy			0.83	90
macro avg	0.83	0.82	0.82	90
weighted avg	0.83	0.83	0.83	90



In [288...

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
#from sklearn.preprocessing import StandardScaler
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
import pandas as pd

df = pd.read_csv(r"C:\Users\vaibh\Downloads\housing.csv")
print("Missing values : ")
print(df.isnull().sum())
df.dropna(inplace = True)
print("\n Duplicate Values : ")
print(df.duplicated().sum())
df.drop_duplicates(inplace = True)
print(df.info())
df = df.select_dtypes(include = 'number')
```

```
X = df.drop(columns=['median_house_value'])
y = df['median_house_value']
print(y)

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.2,random_state =
#scaler = StandardScaler()
#X_train_scaled = scaler.fit_transform(X_train)
#X_test_scaled = scaler.fit_transform(X_test)
model = LinearRegression()
model.fit(X_train,y_train)
y_pred = model.predict(X_test)

mse = mean_squared_error(y_test,y_pred)
print("\nMean Squared Error : ",mse)
mae = mean_absolute_error(y_test,y_pred)
print("Mean absolute square : ",mae)
r2 = r2_score(y_test,y_pred)
print("R2 score : ",r2)
```

```

Missing values :
longitude          0
latitude           0
housing_median_age 0
total_rooms        0
total_bedrooms     207
population         0
households         0
median_income      0
median_house_value 0
ocean_proximity    0
dtype: int64

```

```

Duplicate Values :
0

```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 20433 entries, 0 to 20639
```

```
Data columns (total 10 columns):
```

#	Column	Non-Null Count	Dtype
0	longitude	20433 non-null	float64
1	latitude	20433 non-null	float64
2	housing_median_age	20433 non-null	float64
3	total_rooms	20433 non-null	float64
4	total_bedrooms	20433 non-null	float64
5	population	20433 non-null	float64
6	households	20433 non-null	float64
7	median_income	20433 non-null	float64
8	median_house_value	20433 non-null	float64
9	ocean_proximity	20433 non-null	object

```
dtypes: float64(9), object(1)
```

```
memory usage: 1.7+ MB
```

```
None
```

```

0      452600.0
1      358500.0
2      352100.0
3      341300.0
4      342200.0

```

```
...
```

```

20635    78100.0
20636    77100.0
20637    92300.0
20638    84700.0
20639    89400.0

```

```
Name: median_house_value, Length: 20433, dtype: float64
```

```
Mean Squared Error : 4927778339.313827
```

```
Mean absolute square : 51434.3984174052
```


```
R2 score : 0.6396553423014997
```

In [290...

```

plt.scatter(y_test,y_pred,color = "blue")
plt.title("Actual values vs Predicted Values")
plt.xlabel("Actual Values")
plt.ylabel("Predicted Values")
plt.plot([y_test.min(),y_test.max()], [y_test.min(),y_test.max()],color = "red",line
plt.show()

```

```
In [63]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import accuracy_score, classification_report

df = pd.read_csv(r"C:\Users\vaibh\Downloads\Admission_Predict.csv")
df.head()

df = df.drop(columns = ['Serial No.'])

df['Admit'] = (df['Chance of Admit '] >= 0.75).astype(int)

df = df.drop(columns = ['Chance of Admit '])
X = df.drop(columns = 'Admit')
y = df['Admit']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state

clf = DecisionTreeClassifier()
clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)
print("Accuracy score : ", end = " ")
print(accuracy_score(y_test, y_pred))
```


```
print("Classification Report : ")
print(classification_report(y_test,y_pred))
```

Acucracy score : 0.8375

Classification Report :

	precision	recall	f1-score	support
0	0.79	0.95	0.87	44
1	0.93	0.69	0.79	36
accuracy			0.84	80
macro avg	0.86	0.82	0.83	80
weighted avg	0.85	0.84	0.83	80

```
In [85]: plt.figure(figsize = (10,6))
plot_tree(clf,feature_names=X.columns,class_names= ["Not Adimted","Admitted"])
plt.show()
```



```
In [ ]: #sample code
import sqlite3
import pandas as pd
connection = sqlite3.connect("sample_database.db")
cursor = connection.cursor()

#create table
cursor.execute("""
CREATE TABLE IF NOT EXISTS employees (
    id INTEGER PRIMARY KEY AUTOINCREMENT,
    name TEXT NOT NULL,
    age INTEGER,
    department TEXT,
    salary REAL
```

```

)
"""
cursor.executemany("""
INSERT INTO employees (name,age,department,salary)
VALUES(?,?,?,?)
""", [
    ("Allice",30,"HR",600000),
    ("Bob",25,"IT",70000),
    ("Charlie", 35, "Finance", 80000),
    ("Diana", 28, "IT", 65000),
    ("Eve", 40, "HR", 72000)
])
connection.commit()
connection.close()
print("Data and table setup complete")

connection = sqlite3.connect("sample_database.db")

query = "SELECT * FROM employees"
df = pd.read_sql_query(query,connection)

print("Data loaded into Dataframe : ")
print(df)

connection.close()

print("\n Summary Statistics")
print(df.describe())

df["tax"] = df["salary"]*0.10
print("\n DataFrame with Tax column : ")
print(df)

high_earners = df[df["salary"]>65000]
print("\n Employees with salary > 65000")
print(high_earners)

avg_sal_by_dept = df.groupby("department")["salary"].mean()
print("\nAverage salary by the department : ")
print(avg_sal_by_dept)

sorted_by_age = df.sort_values(by="age",ascending = False)
print("\nEmployees Sorted Age : ")
print(sorted_by_age)

```

```

In [45]: import sqlite3
connection = sqlite3.connect("samp.db")
cursor = connection.cursor()

cursor.execute ("""
CREATE TABLE IF NOT EXISTS employees(
    id INTEGER PRIMARY KEY AUTOINCREMENT,
    name TEXT NOT NULL,
    age INTEGER,

```

```

        department TEXT,
        salary REAL

    )
    """

cursor.executemany("""
INSERT INTO employees(name,age,department,salary)
VALUES(?,?,?,?)
""", [
    ("Alicce",23,"HR",60000),
    ("BOB",34,"IT",70000),
    ("Charlie",25,"Finance",80000),
    ("Dianna",45,"HR",56000),
    ("EVE",40,"HR",72000)
])

connection.commit()
connection.close()
print("Data loaded succesfully")

```

Data loaded succesfully

```

In [47]: import pandas as pd
connection = sqlite3.connect("samp.db")
query = "SELECT * FROM employees"
df = pd.read_sql_query(query,connection)
print("data loaded into dataframe")
df.head()

```

data loaded into dataframe

```

Out[47]:

```

	id	name	age	department	salary
0	1	Alicce	23	HR	60000.0
1	2	BOB	34	IT	70000.0
2	3	Charlie	25	Finance	80000.0
3	4	Dianna	45	HR	56000.0
4	5	EVE	40	HR	72000.0

```

In [59]: df['tax'] = (df['salary']*0.05)
print(df.head())
high_earners = df[df['salary']>65000]
print("\n Emppoyees with salry greter than 65000")
print(high_earners)
avg_sal_by_dept = df.groupby("department").salary.mean()
print("\n Average salary by department")
print(avg_sal_by_dept)
sort_by_age = df.sort_values(by="age",ascending = False)
print("\nsort by age")
print(sort_by_age)

```

	id	name	age	department	salary	tax
0	1	Allice	23	HR	60000.0	3000.0
1	2	BOB	34	IT	70000.0	3500.0
2	3	Charlie	25	Finance	80000.0	4000.0
3	4	Dianna	45	HR	56000.0	2800.0
4	5	EVE	40	HR	72000.0	3600.0

Employees with salry greter than 65000

	id	name	age	department	salary	tax
1	2	BOB	34	IT	70000.0	3500.0
2	3	Charlie	25	Finance	80000.0	4000.0
4	5	EVE	40	HR	72000.0	3600.0

Average salary by department

department

Finance 80000.000000

HR 62666.666667

IT 70000.000000

Name: salary, dtype: float64

sort by age

	id	name	age	department	salary	tax
3	4	Dianna	45	HR	56000.0	2800.0
4	5	EVE	40	HR	72000.0	3600.0
1	2	BOB	34	IT	70000.0	3500.0
2	3	Charlie	25	Finance	80000.0	4000.0
0	1	Allice	23	HR	60000.0	3000.0

```
In [67]: import pandas as pd
import requests
from bs4 import BeautifulSoup

# Fetch the webpage content
url = "https://en.wikipedia.org/wiki/Agriculture_in_India"
response = requests.get(url)

# Parse the HTML using BeautifulSoup
soup = BeautifulSoup(response.text, "html.parser")

# Extract all tables using pandas
tables = pd.read_html(response.text)

# Print all tables
print(f"Number of tables found: {len(tables)}")
for i, table in enumerate(tables):
    print(f"\nTable {i + 1}:")
    print(table.head()) # Display the first few rows of each table
```

	Country or Territory	Population(1 July 2022)	Population(1 July 2023)	\
0	World	8,021,407,192	8,091,734,930	
1	India	1,425,423,212	1,438,069,596	
2	China[a]	1,425,179,569	1,422,584,933	
3	United States	341,534,046	343,477,335	
4	Indonesia	278,830,529	281,190,067	

	Change	UN Continental Region	UN Statistical Subregion
0	+0.88%	–	–
1	+0.89%	Asia	Southern Asia
2	–0.18%	Asia	Eastern Asia
3	+0.57%	Americas	Northern America
4	+0.85%	Asia	South-eastern Asia

```
In [153... import pandas as pd
import requests
from bs4 import BeautifulSoup
from io import StringIO

url = "https://en.wikipedia.org/wiki/Agriculture_in_India"

response = requests.get(url)

soup = BeautifulSoup(response.text,"html.parser")
table = soup.find("table",{"class" : "wikitable"})

df = pd.read_html(StringIO(str(table)))[0]

print("Extracted Data : ")
df.head()
```

Extracted Data :

Out[153...

	Rank	Commodity	Value (US\$, 2016)	Unit price (US\$ / kilogram, 2009)	Average yield (tonnes per hectare, 2017)	Most productive country (tonnes per hectare, 2017)	Most productive country (tonnes per hectare, 2017).1
0	1	Rice	\$70.18 billion	0.27	3.85	9.82	Australia
1	2	Buffalo milk	\$43.09 billion	0.40	2.00[78]	2.00[78]	India
2	3	Cow milk	\$32.55 billion	0.31	1.2[78]	10.3[78]	Israel
3	4	Wheat	\$26.06 billion	0.15	2.8	8.9	Netherlands
4	5	Cotton (Lint + Seeds)	\$23.30 billion	1.43	1.6	4.6	Israel

In [123...

```
# import pandas as pd
# from bs4 import BeautifulSoup
# import requests
# from io import StringIO

# url = "https://en.wikipedia.org/wiki/Agriculture_in_India"
# response = requests.get(url)

# soup = BeautifulSoup(response.text, "html.parser")
# table = soup.find("table", {"class" : "wikitable"})

# df = pd.read_html(StringIO(str(table)))[0]

# rows = table.find_all("tr")
# data = []
# for row in rows:
#     cells = row.find_all(["th", "td"]) # Include both headers and data cells
#     data.append([cell.text.strip() for cell in cells])

# # Convert to DataFrame
# df = pd.DataFrame(data)
# print("Extracted Table Data:")
# print(df.head())
# tables = pd.read_html(response.text)
# for i, table in enumerate(tables):
#     print(f"\nTable {i + 1}:\n", table.head())
```

In [107...

```
df.describe()
```

Out[107...

	0	1	2	3	4	5	6
count	21	21	21	21	21	21	20
unique	21	21	21	19	20	21	16
top	Rank	Commodity	Value (US\$, 2016)	0.4	1.1	Most productive country(tonnes per hectare, 2017)	Israel
freq	1	1	1	2	2	1	3

In [117...

```
#Api integration
import pandas as pd
import requests

url = "https://api.coingecko.com/api/v3/coins/markets"

params = {
    "vs_currency" : "usd",
    "order" : "market_cap_desc",
    "per_page" : 10,
    "page" : 1,
    "sparkLine" : False
}
```

```

response = requests.get(url,params=params)

if response.status_code == 200 :
    data = response.json()
    df = pd.DataFrame(data)
    print("Data succesfully loaded into dataframe")
else :
    print(f"Failed to fetch the data : {response.status_code}")

df.head()

```

Data succesfully loaded into dataframe

Out[117...

	id	symbol	name	image	current_price	ma
0	bitcoin	btc	Bitcoin	https://coin-images.coingecko.com/coins/images...	97041.000000	191756
1	ethereum	eth	Ethereum	https://coin-images.coingecko.com/coins/images...	3107.370000	37371
2	tether	usdt	Tether	https://coin-images.coingecko.com/coins/images...	0.998863	12986
3	solana	sol	Solana	https://coin-images.coingecko.com/coins/images...	239.550000	11355
4	binancecoin	bnb	BNB	https://coin-images.coingecko.com/coins/images...	609.310000	8873

5 rows × 26 columns

In [121...

```

import requests
import pandas as pd

url = "https://api.coingecko.com/api/v3/coins/markets"
params = {
    "vs_currency" : "usd",
    "order" : "market_cap_desc",
    "per_page" : 10,
    "page" : 1,
    "sparkLine" : False
}

response = requests.get(url,params = params)
if response.status_code == 200 :
    data = response.json()
    df = pd.DataFrame(data)
    print("Data succesfully loaded")
else :
    print(f"Failed to load the data {response.status_code}")

```



```
df.head()
```

Data succesfully loaded

Out[121...

	id	symbol	name	image	current_price	ma
0	bitcoin	btc	Bitcoin	https://coin-images.coingecko.com/coins/images...	97041.000000	191756
1	ethereum	eth	Ethereum	https://coin-images.coingecko.com/coins/images...	3107.370000	37377
2	tether	usdt	Tether	https://coin-images.coingecko.com/coins/images...	0.998863	12986
3	solana	sol	Solana	https://coin-images.coingecko.com/coins/images...	239.550000	11355
4	binancecoin	bnb	BNB	https://coin-images.coingecko.com/coins/images...	609.310000	8875

5 rows × 26 columns



In [139...

```
import pandas as pd

df = pd.read_csv(r"C:\Users\vaibh\Downloads\heart (1).csv")
print(df)
print("\nMissing values : ")
print(df.isnull().sum())

df_dropped = df.dropna()
print("\n Data after dropping Missing values : ")
print(df_dropped)

df_specific_value = df.fillna(0)
print("\nFilling with Specific Value(0) ")
print(df_specific_value)

df_fill_mean = df.fillna(df.mean())
print("Filling missing values with mean: ")
print(df_fill_mean)
```

	age	gender	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63.0	1.0	3	145.0	233	1	0	150	0	2.3	
1	NaN	NaN	2	130.0	250	0	1	187	0	3.5	
2	41.0	0.0	1	130.0	204	0	0	172	0	1.4	
3	NaN	1.0	1	120.0	236	0	1	178	0	0.8	
4	57.0	NaN	0	120.0	354	0	1	163	1	0.6	
..	
298	57.0	0.0	0	140.0	241	0	1	123	1	0.2	
299	45.0	1.0	3	110.0	264	0	1	132	0	1.2	
300	68.0	1.0	0	144.0	193	1	1	141	0	3.4	
301	57.0	1.0	0	130.0	131	0	1	115	1	1.2	
302	57.0	0.0	1	130.0	236	0	0	174	0	0.0	

	slope	ca	thal	target
0	0	0	1	1
1	0	0	2	1
2	2	0	2	1
3	2	0	2	1
4	2	0	2	1
..
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

[303 rows x 14 columns]

Missing values :

```

age      3
gender    2
cp        0
trestbps  1
chol      0
fbs       0
restecg   0
thalach   0
exang     0
oldpeak   0
slope     0
ca        0
thal      0
target    0
dtype: int64

```

Data after dropping Missing values :

	age	gender	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63.0	1.0	3	145.0	233	1	0	150	0	2.3	
2	41.0	0.0	1	130.0	204	0	0	172	0	1.4	
5	57.0	1.0	0	140.0	192	0	1	148	0	0.4	
7	44.0	1.0	1	120.0	263	0	1	173	0	0.0	
8	52.0	1.0	2	172.0	199	1	1	162	0	0.5	
..	
298	57.0	0.0	0	140.0	241	0	1	123	1	0.2	
299	45.0	1.0	3	110.0	264	0	1	132	0	1.2	
300	68.0	1.0	0	144.0	193	1	1	141	0	3.4	

301	57.0	1.0	0	130.0	131	0	1	115	1	1.2
302	57.0	0.0	1	130.0	236	0	0	174	0	0.0

	slope	ca	thal	target
0	0	0	1	1
2	2	0	2	1
5	1	0	1	1
7	2	0	3	1
8	2	0	3	1
..
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

[299 rows x 14 columns]

Filling with Specific Value(0)

	age	gender	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63.0	1.0	3	145.0	233	1	0	150	0	2.3	
1	0.0	0.0	2	130.0	250	0	1	187	0	3.5	
2	41.0	0.0	1	130.0	204	0	0	172	0	1.4	
3	0.0	1.0	1	120.0	236	0	1	178	0	0.8	
4	57.0	0.0	0	120.0	354	0	1	163	1	0.6	
..	
298	57.0	0.0	0	140.0	241	0	1	123	1	0.2	
299	45.0	1.0	3	110.0	264	0	1	132	0	1.2	
300	68.0	1.0	0	144.0	193	1	1	141	0	3.4	
301	57.0	1.0	0	130.0	131	0	1	115	1	1.2	
302	57.0	0.0	1	130.0	236	0	0	174	0	0.0	

	slope	ca	thal	target
0	0	0	1	1
1	0	0	2	1
2	2	0	2	1
3	2	0	2	1
4	2	0	2	1
..
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

[303 rows x 14 columns]

Filling missing values with mean:

	age	gender	cp	trestbps	chol	fbs	restecg	thalach	exang	\
0	63.000000	1.000000	3	145.0	233	1	0	150	0	
1	54.413333	0.684385	2	130.0	250	0	1	187	0	
2	41.000000	0.000000	1	130.0	204	0	0	172	0	
3	54.413333	1.000000	1	120.0	236	0	1	178	0	
4	57.000000	0.684385	0	120.0	354	0	1	163	1	
..	
298	57.000000	0.000000	0	140.0	241	0	1	123	1	
299	45.000000	1.000000	3	110.0	264	0	1	132	0	

300	68.000000	1.000000	0	144.0	193	1	1	141	0
301	57.000000	1.000000	0	130.0	131	0	1	115	1
302	57.000000	0.000000	1	130.0	236	0	0	174	0

	oldpeak	slope	ca	thal	target
0	2.3	0	0	1	1
1	3.5	0	0	2	1
2	1.4	2	0	2	1
3	0.8	2	0	2	1
4	0.6	2	0	2	1
..
298	0.2	1	0	3	0
299	1.2	1	0	3	0
300	3.4	1	2	3	0
301	1.2	1	1	3	0
302	0.0	1	1	2	0

[303 rows x 14 columns]

In [151...

```
import pandas as pd
from scipy.stats import zscore
import numpy as np
df = pd.read_csv(r"C:\Users\vaibh\Downloads\heart (1).csv")
z_score = zscore(df,nan_policy = "omit")
print(abs_z_score)
abs_z_score = np.abs(z_score)
threshold = 2

df_cleaned = df.loc[(abs_z_score < threshold).all(axis = 1)]
print("\nData cleaned : ")
print(df_cleaned)
```

	age	gender	cp	trestbps	chol	fbs	restecg \
0	0.948186	0.679091	1.973123	0.764565	0.256334	2.394438	1.005832
1	NaN	NaN	1.002577	0.091038	0.072199	0.417635	0.898962
2	1.481172	1.472556	0.032031	0.091038	0.816773	0.417635	1.005832
3	NaN	0.679091	0.032031	0.661440	0.198357	0.417635	0.898962
4	0.285634	NaN	0.938515	0.661440	2.082050	0.417635	0.898962
..
298	0.285634	1.472556	0.938515	0.479364	0.101730	0.417635	0.898962
299	1.039471	0.679091	1.973123	1.231842	0.342756	0.417635	0.898962
300	1.500313	0.679091	0.938515	0.707525	1.029353	2.394438	0.898962
301	0.285634	0.679091	0.938515	0.091038	2.227533	0.417635	0.898962
302	0.285634	1.472556	0.032031	0.091038	0.198357	0.417635	1.005832

	thalach	exang	oldpeak	slope	ca	thal	target
0	0.015443	0.696631	1.087338	2.274579	0.714429	2.148873	0.914529
1	1.633471	0.696631	2.122573	2.274579	0.714429	0.512922	0.914529
2	0.977514	0.696631	0.310912	0.976352	0.714429	0.512922	0.914529
3	1.239897	0.696631	0.206705	0.976352	0.714429	0.512922	0.914529
4	0.583939	1.435481	0.379244	0.976352	0.714429	0.512922	0.914529
..
298	1.165281	1.435481	0.724323	0.649113	0.714429	1.123029	1.093459
299	0.771706	0.696631	0.138373	0.649113	0.714429	1.123029	1.093459
300	0.378132	0.696631	2.036303	0.649113	1.244593	1.123029	1.093459
301	1.515125	1.435481	0.138373	0.649113	0.265082	1.123029	1.093459
302	1.064975	0.696631	0.896862	0.649113	0.265082	0.512922	1.093459

[303 rows x 14 columns]

Data cleaned :

	age	gender	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak \
2	41.0	0.0	1	130.0	204	0	0	172	0	1.4
7	44.0	1.0	1	120.0	263	0	1	173	0	0.0
9	57.0	1.0	2	150.0	168	0	1	174	0	1.6
10	54.0	1.0	0	140.0	239	0	1	160	0	1.2
11	48.0	0.0	2	130.0	275	0	1	139	0	0.2
..
293	67.0	1.0	2	152.0	212	0	0	150	0	0.8
296	63.0	0.0	0	124.0	197	0	1	136	1	0.0
298	57.0	0.0	0	140.0	241	0	1	123	1	0.2
299	45.0	1.0	3	110.0	264	0	1	132	0	1.2
302	57.0	0.0	1	130.0	236	0	0	174	0	0.0

	slope	ca	thal	target
2	2	0	2	1
7	2	0	3	1
9	2	0	2	1
10	2	0	2	1
11	2	0	2	1
..
293	1	0	3	0
296	1	0	2	0
298	1	0	3	0
299	1	0	3	0
302	1	1	2	0

[176 rows x 14 columns]

```
In [49]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv(r"C:\Users\vaibh\Downloads\housing.csv")
print("Missing values : ")
print(df.isnull().sum())
df.dropna(inplace = True)
print("Duplicate values : ")
print(df.duplicated().sum())
df.drop_duplicates(inplace = True)

df = df.select_dtypes(include = "number")
X = df.drop(columns = ['median_house_value'])
y = df['median_house_value']
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.3,random_state =
model = LinearRegression()
model.fit(X_train,y_train)

y_pred = model.predict(X_test)

mse = mean_squared_error(y_test,y_pred)
print("\nMean squared Error",end = " ")
print(mse)
print("Mean absolute square : ",mean_absolute_error(y_test,y_pred))
print("r2 score : ",r2_score(y_test,y_pred))

plt.figure(figsize = (10,6))
plt.scatter(y_test,y_pred,color = "blue")
plt.title("Actual Value vs Predicted Values")
plt.xlabel("Actual values")
plt.ylabel("Predicted values")
plt.plot([y_test.min(),y_test.max()], [y_test.min(),y_test.max()],color = "red",line
plt.show()
```

```


Missing values :
longitude          0
latitude           0
housing_median_age  0
total_rooms        0
total_bedrooms     207
population         0
households         0
median_income      0
median_house_value  0
ocean_proximity    0
dtype: int64
Duplicate values :
0

```

```

Mean squared Error 4738972791.400478
Mean absolute square : 50704.919692847856
r2 score : 0.6445130291082346

```



```

In [45]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
import pandas as pd

df = pd.read_csv(r"C:\Users\vaibh\Downloads\heart (1).csv")
print("Missing values : ")
print(df.isnull().sum())
df.dropna(inplace = True)
print("\nDuplicate values : ")
print(df.duplicated().sum())
df.drop_duplicates(inplace = True)

X = df.iloc[:,0:13]
y = df['target']

```

```

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.3,random_state =
model = LogisticRegression(max_iter = 1000)
model.fit(X_train,y_train)

y_pred = model.predict(X_test)
classification = classification_report(y_test,y_pred)
confusion = confusion_matrix(y_test,y_pred)
accuracy = accuracy_score(y_test,y_pred)

print("\nAccuracy Score : ")
print(accuracy)
print("\nClassification Report : ")
print(classification)
print("\nconfusion Matrix : ")
print(confusion)

```

Missing values :

```

age      3
gender   2
cp        0
trestbps 1
chol      0
fbs       0
restecg   0
thalach   0
exang     0
oldpeak   0
slope     0
ca        0
thal      0
target    0
dtype: int64

```

Duplicate values :

```
1
```

Accuracy Score :

```
0.8333333333333334
```

Classification Report :

	precision	recall	f1-score	support
0	0.82	0.75	0.78	36
1	0.84	0.89	0.86	54
accuracy			0.83	90
macro avg	0.83	0.82	0.82	90
weighted avg	0.83	0.83	0.83	90

confusion Matrix :

```

[[27  9]
 [ 6 48]]

```



```
In [92]: import pandas as pd
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt

df = pd.read_csv(r"C:\Users\vaibh\Downloads\Admission_Predict.csv")
df.dropna(inplace = True)
df.drop_duplicates(inplace = True)
df = df.select_dtypes(include = "number")
df = df.drop(columns = ['Serial No.'])
df['Admit'] = (df['Chance of Admit '] >= 0.75).astype(int)
df = df.drop(columns = ['Chance of Admit '])
X = df.drop(columns = ['Admit'])
y = df['Admit']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 42)
clf = DecisionTreeClassifier()
clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)
classification = classification_report(y_test, y_pred)
confusion = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix")
print(confusion)
print("\nClassification Report")
print(classification)
```

Confusion Matrix

```
[[59  7]
 [13 41]]
```

Classification Report

	precision	recall	f1-score	support
0	0.82	0.89	0.86	66
1	0.85	0.76	0.80	54
accuracy			0.83	120
macro avg	0.84	0.83	0.83	120
weighted avg	0.84	0.83	0.83	120

```
In [98]: plt.figure(figsize = (10,6))
plot_tree(clf, feature_names = X.columns, class_names = ['Not Admitted', 'Admitted'])
plt.show()
```

```
In [25]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
df = pd.read_csv(r"C:\Users\vaibh\Downloads\iris.csv")
print("missing values : ")
print(df.isnull().sum())
print("Duplicates : ")
print(df.duplicated().sum())
df.dropna(inplace = True)
df.drop_duplicates(inplace = True)
corelation_matrix = df.select_dtypes(include = "number").corr()
plt.figure(figsize=(10,6))
```

```
sns.heatmap(corelation_matrix,annot = True,cmap = 'coolwarm',fmt = '0.2f',cbar = Tr
plt.xticks(rotation = 45)
plt.yticks(rotation = 45)
plt.show()
```

missing values :

sepal.length 0

sepal.width 0

petal.length 0


petal.width 0

species 0

dtype: int64

Duplicates :

1



In []: